

Blockchains & Distributed Ledgers

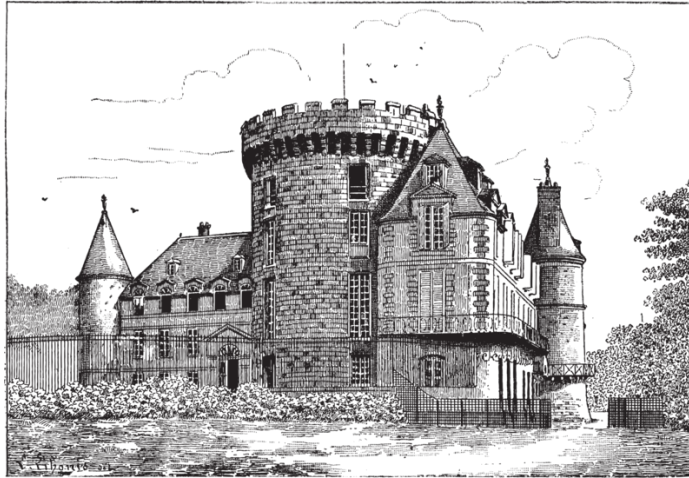
Lecture 05

Dimitris Karakostas



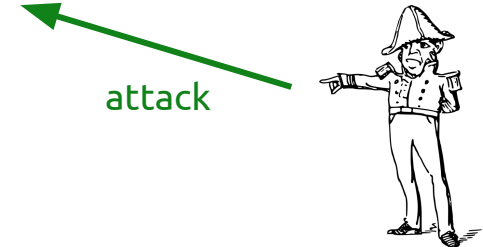
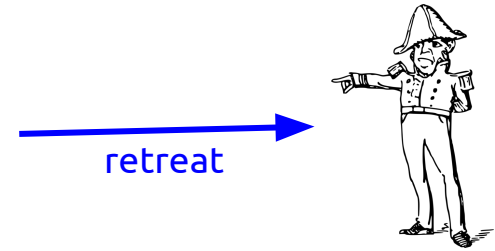
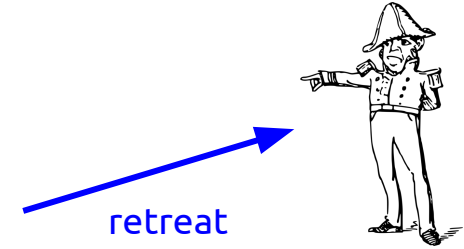
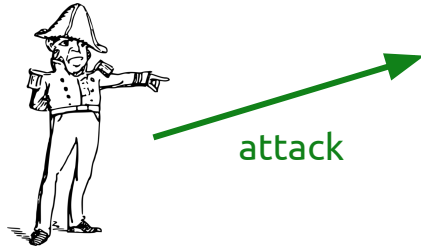
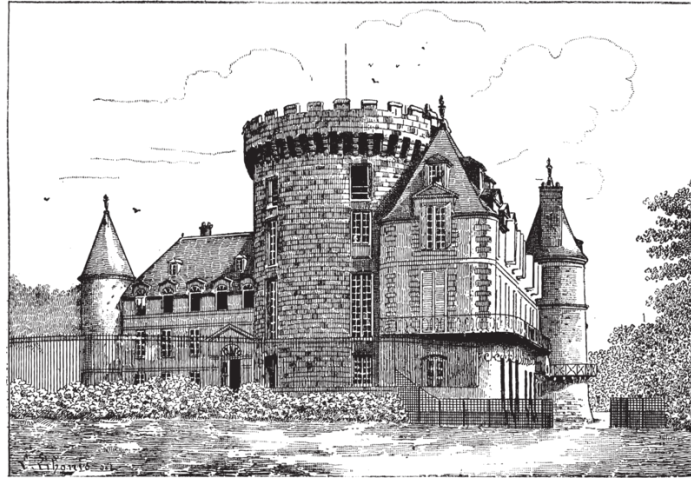

Slide credits: DK, Aggelos Kiayias, Dionysis Zindros, Christos Nasikas

The Byzantine Generals Problem

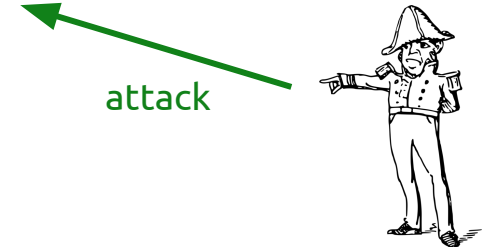
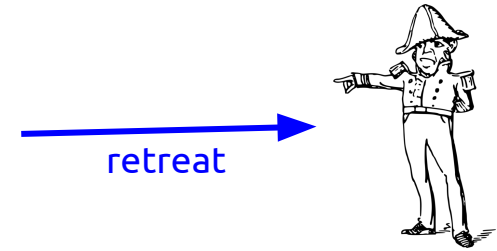
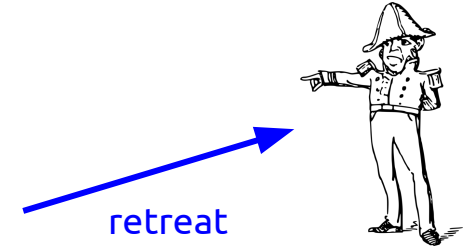
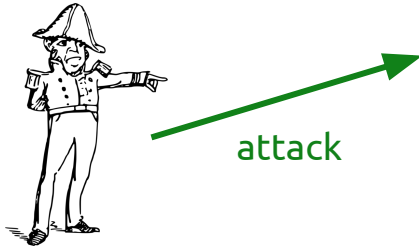

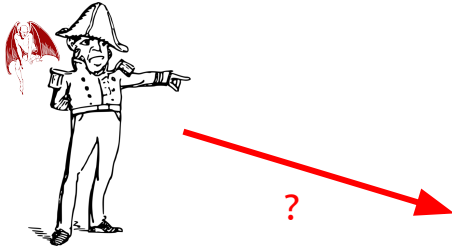


[LSP82] Leslie Lamport, Robert Shostak, and Marshall Pease.
The byzantine generals problem

The Byzantine Generals Problem



The Byzantine Generals Problem



The Consensus Problem

Motivation for the Consensus Layer, I

- A transaction history and/or state of the service needs to be **agreed** by all servers.
- Servers may be operated by participants with **diverging interests**, in terms of the history of transactions and/or state of the service.

Motivation for the Consensus Layer, II

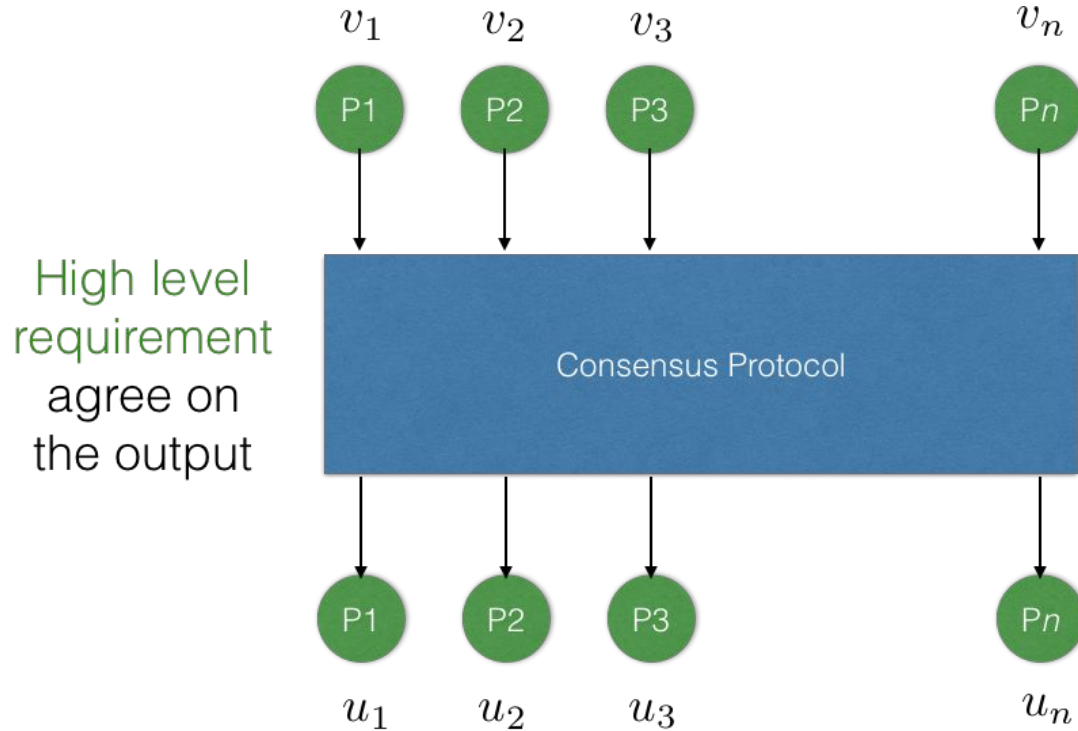


Consensus : Problem Statement

- A number (t) of the participating entities can diverge from the protocol.
- This has been called **Byzantine behaviour** in the literature.
- The properties of the protocol are defined in the presence of this “malicious” coalition of parties that attempts to disrupt the process for the “honest” parties.

$$H, |H| = n - t$$

The consensus problem



Study initiated by Lamport, Pease, Shostak 1982

Consensus Properties

- Termination $\forall i \in H(u_i \text{ is defined})$

Consensus Properties

- Termination $\forall i \in \mathbf{H}(u_i \text{ is defined})$
- Agreement $\forall i, j \in \mathbf{H}(u_i = u_j)$

Consensus Properties

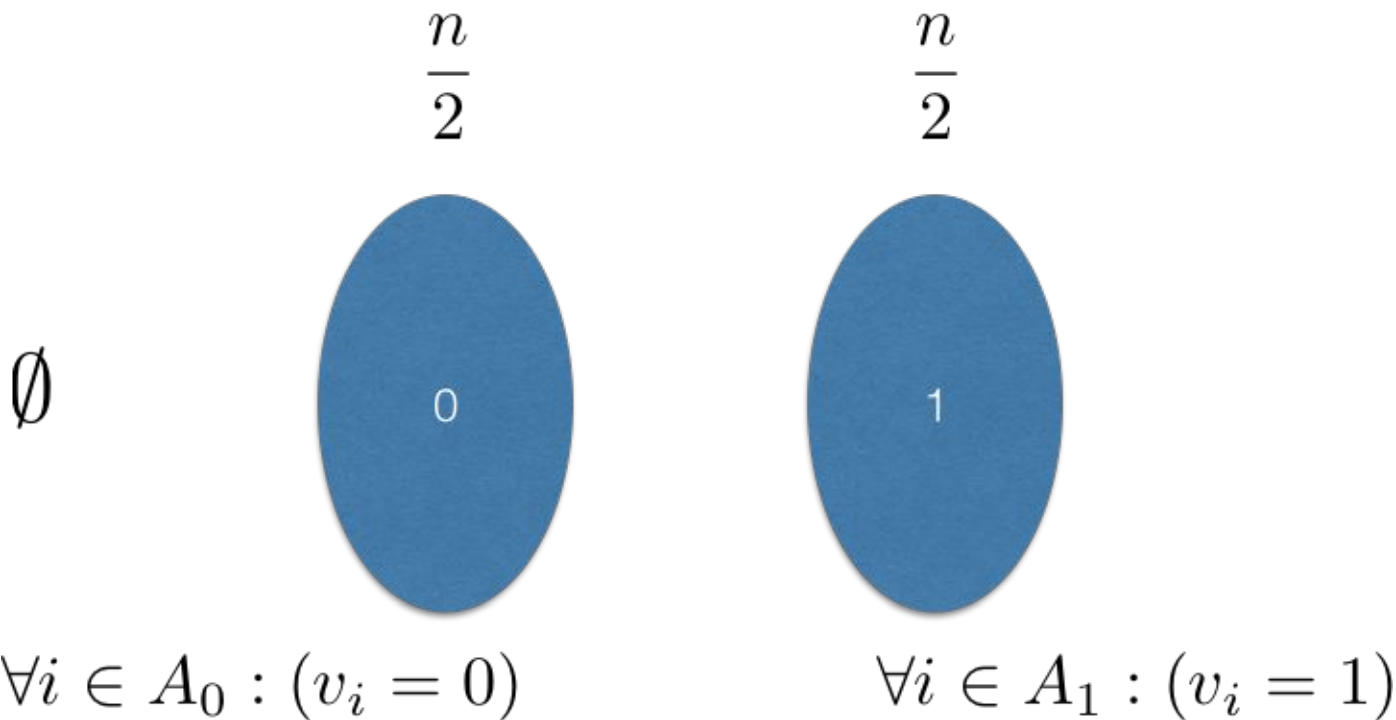
- Termination $\forall i \in \mathbf{H}(u_i \text{ is defined})$
- Agreement $\forall i, j \in \mathbf{H}(u_i = u_j)$
- Validity $\exists v(\forall i \in \mathbf{H}(v_i = v)) \implies (\forall i \in \mathbf{H}(u_i = v))$

Consensus Properties

- Termination $\forall i \in H (u_i \text{ is defined})$
- Agreement $\forall i, j \in H (u_i = u_j)$
- Validity $\exists v (\forall i \in H (v_i = v)) \implies (\forall i \in H (u_i = v))$
- Strong Validity $\forall i \in H \exists j \in H (u_i = v_j)$

Honest Majority is Necessary, I

Consider an adversary that performs one of the following with probability $\frac{1}{3}$



Honest Majority is Necessary, II

- If consensus protocol secure:
 - Adversary corrupts A_0 : output of honest parties (that belong to A_1) should be 1.
 - Adversary corrupts A_1 : output of honest parties (that belong to A_0) should be 0.
 - Adversary corrupts no-one: output of all parties should be the same.
- Adversary corrupts each set with prob. $\frac{1}{3}$ and instructs corrupted parties to follow the protocol
 - honest parties cannot distinguish between honest/corrupted parties
- If all parties output same value: validity is violated with prob. at least $\frac{1}{3}$
- If all parties output different value: consistency is violated with prob. at least $\frac{1}{3}$

Is Honest Majority Sufficient?

- Two important scenarios have been considered in the consensus literature.
 - Point to point channels. **No setup.**
 - Point to point channels. **With setup.**
- The setup provides a correlated private initialization string to each participant; *it is assumed* to be honestly produced.

Setup and Network

Setup/Network	Synchrony	Partial Synchrony
No Setup	$t < n/3$	$t < n/3$
With Setup	$t < n/2$	$t < n/3$

We know consensus can be achieved, assuming the above bounds on adversarial parties.

The typical setup and network configuration in classical consensus protocols

- Setup: a public-key directory
 - Parties have signing and verification keys for a digital signature scheme.
 - Each party knows every other party's verification key.
- Network: point-to-point channels
 - Synchronous, partially synchronous, or asynchronous

Bitcoin Consensus

Enter Bitcoin (2008-09)

- Important concepts used by Bitcoin
 - blockchain data structure
 - proof of work (POW)
- Both known and studied earlier, but combined for a novel application

The setup and network configuration in Bitcoin

- Setup: a random (unpredictable) string
 - The blockchain protocol runs without relying on public-key crypto
- Network: peer-to-peer diffusion
 - Synchronous for at least a small subset of the participants (that may be evolving over time).

The Bitcoin Setting for Consensus

- Also referred to as the “permissionless” setting.
- The bitcoin setting is different, compared to what has been considered classically for the consensus problem.
- Communication is by **diffusion** (no point-to-point channels).
 - Message delivery is assumed, but message origins and recipient list are not specified.
- The protocol setup is **not** a private correlated setup
 - Digital signatures are **not** used to authenticate miners
 - A public setup is assumed: a genesis block

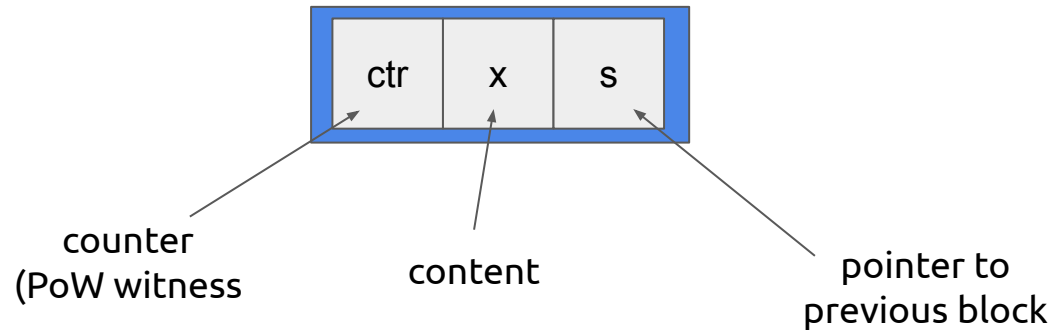
The Bitcoin “backbone”

- The core of the bitcoin protocol
 - The chain validation predicate.
 - The chain selection rule (max-valid)
 - The proof of work function.
 - The main protocol loop
- Protocol is executed by “miners”

[GKL2015] Garay, Kiayias, Leonardos. The Bitcoin Backbone Protocol: Analysis and Applications.

Model

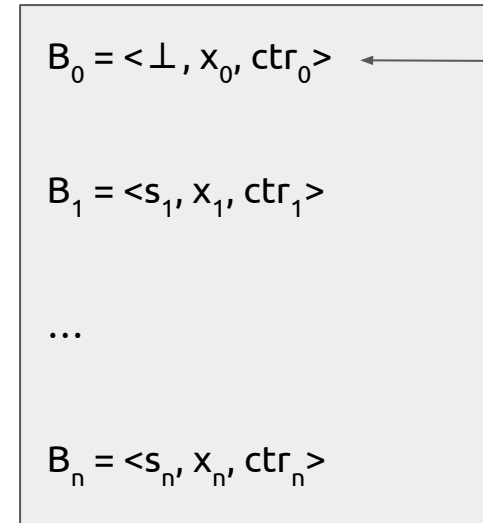
- Assume there are n parties running the protocol
- Synchronous
- Each party has a quota of q queries to the function $H(\cdot)$ in each round
- A number of t parties are controlled by an adversary (a malicious coalition)
 - Security arguments are *for any adversary*



Algorithm 3 The *proof of work* function, parameterized by q , T and hash functions $H(\cdot), G(\cdot)$.
The input is (x, \mathcal{C}) .

```
1: function pow( $x, \mathcal{C}$ )
2:   if  $\mathcal{C} = \varepsilon$  then                                      $\triangleright$  Determine proof of work instance
3:      $s \leftarrow 0$ 
4:   else
5:      $\langle s', x', ctr' \rangle \leftarrow \text{head}(\mathcal{C})$ 
6:      $s \leftarrow H(ctr', G(s', x'))$ 
7:   end if
8:    $ctr \leftarrow 1$ 
9:    $B \leftarrow \varepsilon$ 
10:   $h \leftarrow G(s, x)$ 
11:  while ( $ctr \leq q$ ) do
12:    if ( $H(ctr, h) < T$ ) then                                $\triangleright$  This  $H(\cdot)$  invocation subject to the  $q$ -bound
13:       $B \leftarrow \langle s, x, ctr \rangle$ 
14:      break
15:    end if
16:     $ctr \leftarrow ctr + 1$ 
17:  end while
18:   $\mathcal{C} \leftarrow \mathcal{C}B$                                         $\triangleright$  Extend chain
19:  return  $\mathcal{C}$ 
20: end function
```

Blockchain



$$B_0 = \langle \perp, x_0, \text{ctr}_0 \rangle$$

genesis block

$$B_1 = \langle s_1, x_1, \text{ctr}_1 \rangle$$

...

$$B_n = \langle s_n, x_n, \text{ctr}_n \rangle$$

$$s_i = H(\text{ctr}_{i-1}, G(x_{i-1}, s_{i-1}))$$

$$x_C = \langle x_0, x_1, \dots, x_n \rangle$$

$$C^{\text{rk}} = \langle B_0, B_1, \dots, B_{n-k} \rangle$$

$$C = \langle B_0, B_1, \dots, B_n \rangle$$

chain

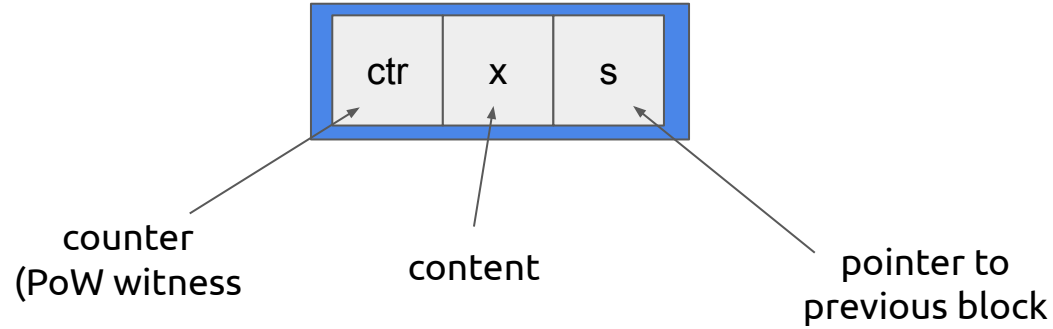
head

Algorithm 1 The *chain validation predicate*, parameterized by q, T , the hash functions $G(\cdot), H(\cdot)$, and the *content validation predicate* $V(\cdot)$. The input is \mathcal{C} .

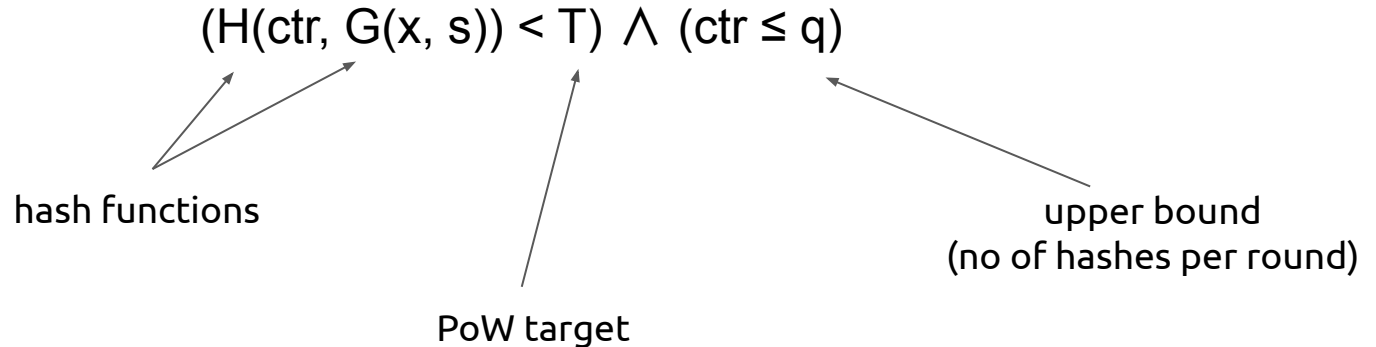
```

1: function validate( $\mathcal{C}$ )
2:    $b \leftarrow V(\mathbf{x}_{\mathcal{C}})$ 
3:   if  $b \wedge (\mathcal{C} \neq \varepsilon)$  then                                      $\triangleright$  The chain is non-empty and meaningful w.r.t.  $V(\cdot)$ 
4:      $\langle s, x, ctr \rangle \leftarrow \text{head}(\mathcal{C})$ 
5:      $s' \leftarrow H(ctr, G(s, x))$ 
6:     repeat
7:        $\langle s, x, ctr \rangle \leftarrow \text{head}(\mathcal{C})$ 
8:       if  $\text{validblock}_q^T(\langle s, x, ctr \rangle) \wedge (H(ctr, G(s, x)) = s')$  then
9:          $s' \leftarrow s$                                               $\triangleright$  Retain hash value
10:         $\mathcal{C} \leftarrow \mathcal{C}^{\uparrow 1}$                                       $\triangleright$  Remove the head from  $\mathcal{C}$ 
11:       else
12:          $b \leftarrow \text{False}$ 
13:       end if
14:     until  $(\mathcal{C} = \varepsilon) \vee (b = \text{False})$ 
15:   end if
16:   return ( $b$ )
17: end function

```



validblock predicate:



Algorithm 2 The function that finds the “best” chain, parameterized by function $\max(\cdot)$. The input is $\{\mathcal{C}_1, \dots, \mathcal{C}_k\}$.

```
1: function maxvalid( $\mathcal{C}_1, \dots, \mathcal{C}_k$ )
2:    $temp \leftarrow \varepsilon$ 
3:   for  $i = 1$  to  $k$  do
4:     if validate( $\mathcal{C}_i$ ) then
5:        $temp \leftarrow \max(\mathcal{C}_i, temp)$ 
6:     end if
7:   end for
8:   return  $temp$ 
9: end function
```

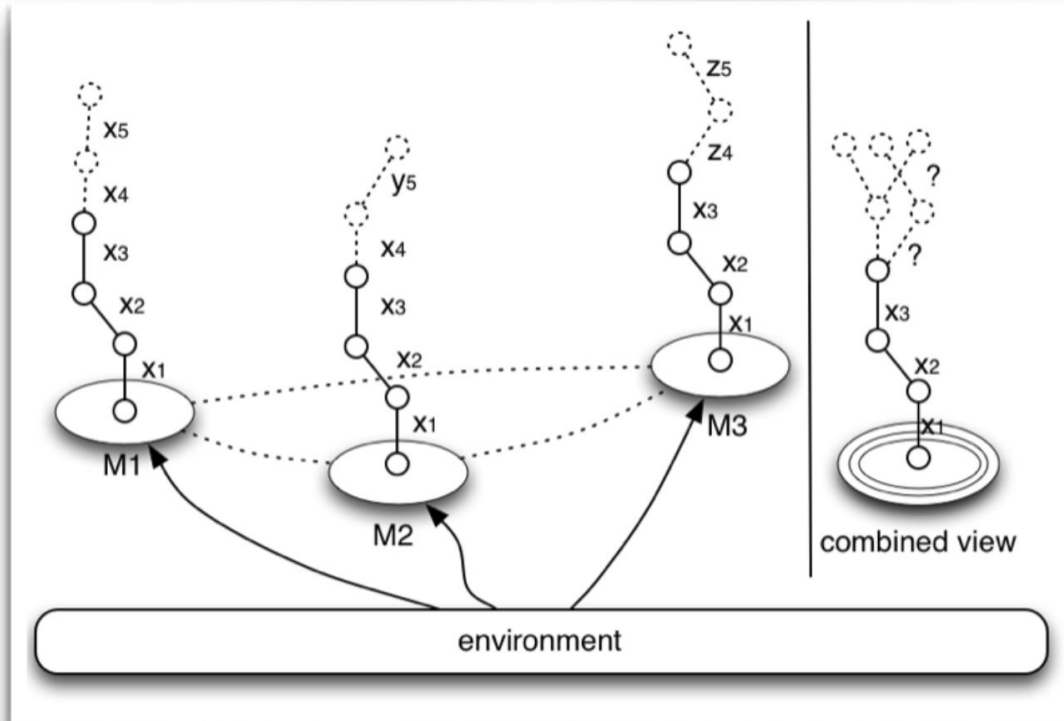
Algorithm 4 The Bitcoin backbone protocol, parameterized by the *input contribution function* $I(\cdot)$ and the *chain reading function* $R(\cdot)$. At the onset it is assumed “init= True”.

```
1: if (init) then
2:    $\mathcal{C} \leftarrow \varepsilon$ 
3:    $st \leftarrow \varepsilon$ 
4:    $round \leftarrow 1$ 
5:    $init \leftarrow \text{False}$ 
6: else
7:    $\tilde{\mathcal{C}} \leftarrow \text{maxvalid}(\mathcal{C}, \text{any chain } \mathcal{C}' \text{ found in RECEIVE}())$ 
8:   if INPUT() contains READ then
9:     write  $R(\tilde{\mathcal{C}})$  to OUTPUT()  $\triangleright$  Produce necessary output before the POW stage.
10:  end if
11:   $\langle st, x \rangle \leftarrow I(st, \tilde{\mathcal{C}}, round, \text{INPUT}(), \text{RECEIVE}())$   $\triangleright$  Determine the  $x$ -value.
12:   $\mathcal{C}_{\text{new}} \leftarrow \text{pow}(x, \tilde{\mathcal{C}})$ 
13:  if  $\mathcal{C} \neq \mathcal{C}_{\text{new}}$  then
14:     $\mathcal{C} \leftarrow \mathcal{C}_{\text{new}}$ 
15:    DIFFUSE( $\mathcal{C}$ )  $\triangleright$  Broadcast the chain in case of adoption/extension.
16:  else
17:    DIFFUSE( $\perp$ )  $\triangleright$  Signals the end of the round to the diffuse functionality.
18:  end if
19:   $round \leftarrow round + 1$ 
20: end if
```

Basic Properties

- Common Prefix
- Chain Quality
- Chain Growth

Common Prefix, I



Common Prefix, II

(strong common prefix / consistency)

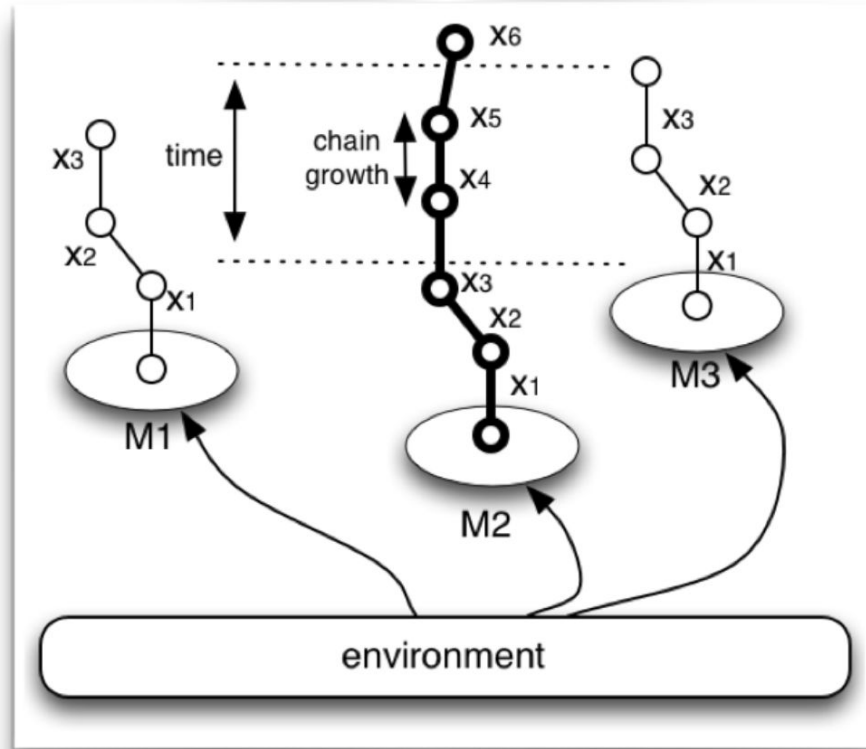
$$\forall r_1, r_2, (r_1 \leq r_2), P_1, P_2, \text{ with } \mathcal{C}_1, \mathcal{C}_2 : \mathcal{C}_1^{\lceil k} \preceq \mathcal{C}_2$$

- The property holds true, in a probabilistic sense, with an error that decays exponentially in k

Racing Attacks

- Attacker splits from the main chain and tries to overtake the “honest chain”
=> Common prefix breaks
- Intuition why the attack is a small probability event:
concentration bounds help honest parties

Chain Growth, I



Chain Growth, II

Parameters $\tau \in (0, 1)$, $s \in \mathbb{N}$

In any period of s rounds at least τs blocks are added to the chain of an honest party P .

- The property holds true in a probabilistic sense with an error probability that exponentially decays in s

$\tau \approx$ probability at least one honest party finds a POW in a round

Abstention Attacks

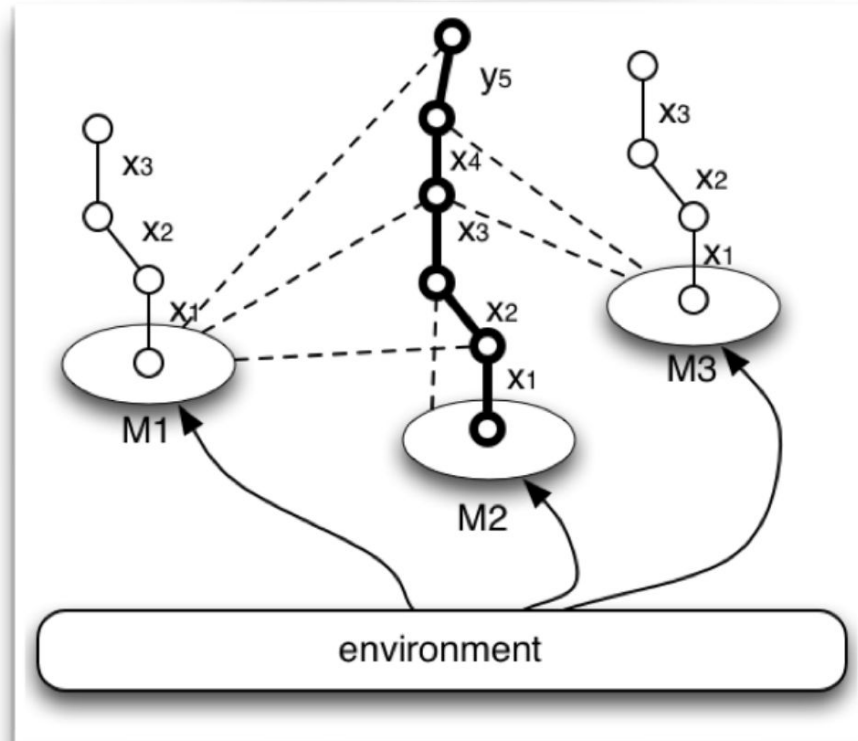
- Attacker stops producing blocks

=> Chain growth stops

- Intuition why the attack is a small probability event:

honest parties will eventually issue blocks

Chain Quality, I



Chain Quality, II

Parameters $\mu \in \{0, 1\}, \ell \in \mathbb{N}$

The ratio of blocks of an ℓ -long segment of an honest chain produced by the adversary is bounded by $(1 - \mu)\ell$

- The property holds true probabilistically with an error that exponentially decays in ℓ

$$\mu \approx \frac{n - 2t}{n - t}$$

Block Withholding Attacks

- Attacker mines privately and releases their block at the same time an honest party releases its own block
- Assuming honest propagation favours the adversary, the honest block is dropped, reducing chain quality
- Intuition why the attack is a small probability event:
over time the adversary cannot produce blocks at the same rate as honest parties (to compete with them)

Robust Transaction Ledger (RTL) - Ledger Consensus

- It can be shown that the three properties can provide a ledger with two core characteristics
- **Persistence:** Transactions are organized in a “*log*” and honest nodes agree on it.
- **Liveness:** New transactions are included in the log, after a suitable (upper-bounded) period of time.

Establishing a RTL from a Blockchain

- Persistence \leftarrow (*strong*) *Common Prefix*
 - need to exclude k most recent blocks
- Liveness \leftarrow *Chain Growth* and *Chain Quality*
 - leave sufficient time for chain to grow
 - apply chain quality to ensure that at least one honest block is included

Ledger Consensus vs. Consensus

- What is the connection?
 - ledger is an ever-going protocol with inputs (e.g., transactions) continuously coming from also external sources
 - consensus is a one-shot execution
- Is it possible to reduce consensus to the ledger? Is it possible to reduce the ledger to consensus?
 - (See the [GKL paper](#) for more details)

Hash operations

- Consider a regular PC (30 MHash / sec)
- With expectation of 2^{74} hashing operations, mining a block will require ~ 20 *million* years.

https://en.bitcoin.it/wiki/Non-specialized_hardware_comparison

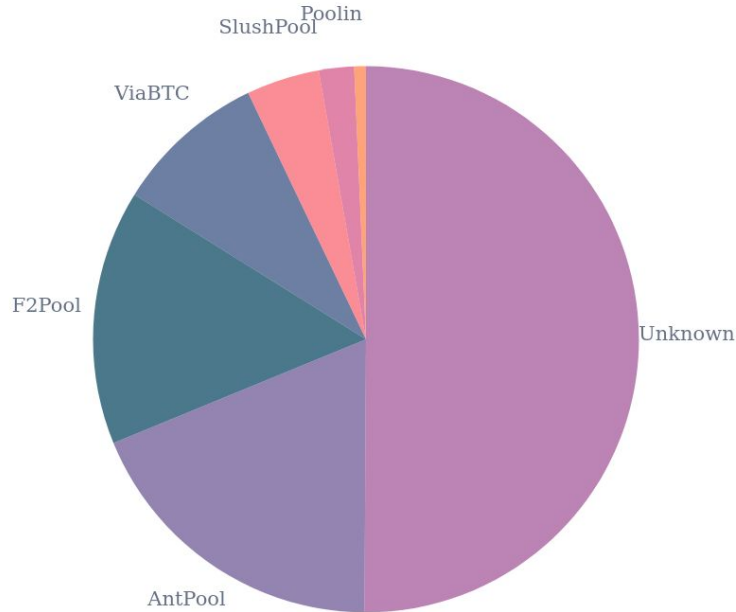
Parallelising mining

- Bitcoin's Proof of Work can be parallelized
- Parties tend to form mining *pools*
 - Instead of working separately, work **together** to solve PoW for the same block.
 - By collecting “**shares**” (small hashes of the block that are not quite as small as needed) one can prove how much they contributed.

Bitcoin mining pools

Hashrate Distribution

An estimation of hashrate distribution amongst the largest mining pools.



<https://www.blockchain.com/pools>

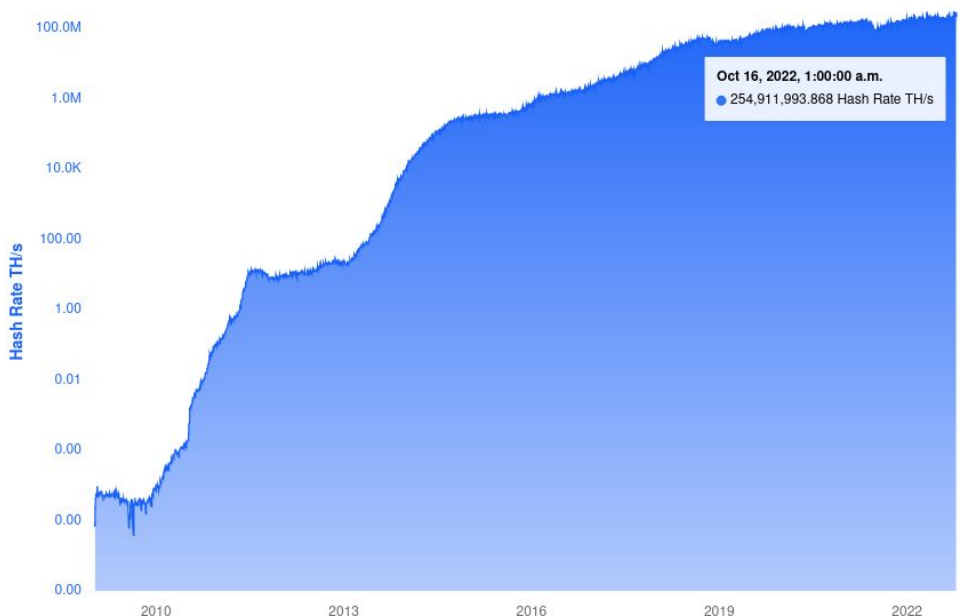
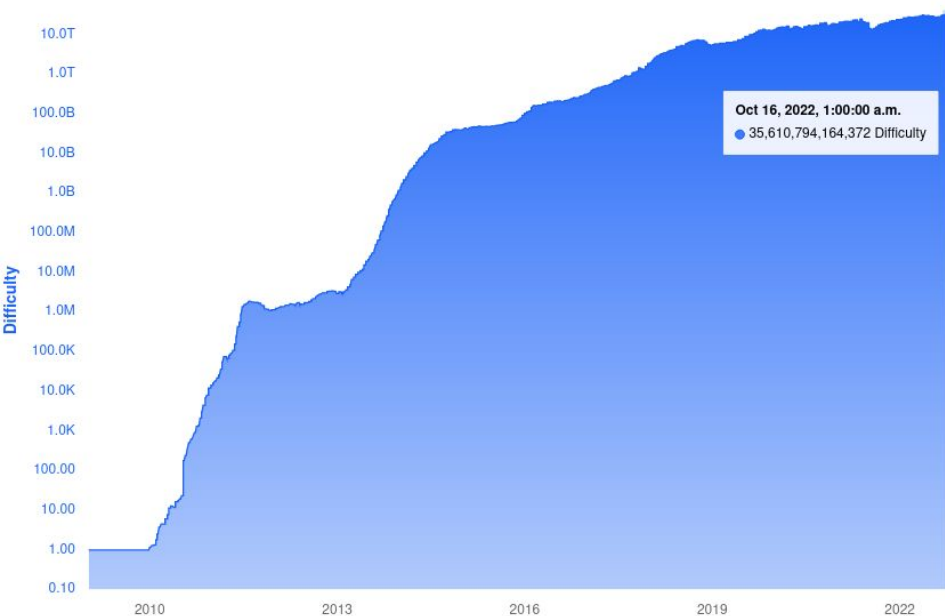
Recall: PoW algorithm

```
int counter;  
counter = 0  
while Hash(data, counter) > Target  
    increment counter  
return counter
```


Dynamic Availability

- So far: n nodes maintain the blockchain
- This number may change over time:
 - new users enter the system
 - existing users leave
- The change over time can be dramatic
- The Bitcoin blockchain handles this, by adjusting the target (difficulty) of the Proof of Work algorithm

Target difficulty / Total hash rate over time



Adjusting the difficulty

“maxvalid” rule is changed

s.t. parties adopt **chain with highest difficulty** linearly related to:

$$\sum_i \frac{1}{T_i}$$

The f parameter [GKL15]

f = probability of producing a block in a round of interaction

- f depends on:
 - target T
 - number of miners
 - duration of round
- If f becomes too small, parties do not progress
 - Chain growth slows
 - Liveness is hurt
- If f becomes too large, parties “collide” often
 - Attacker can exploit network scheduling of message delivery to create forks
 - Persistence is hurt
- To resolve this dynamically, Bitcoin **recalculates** T to keep f constant

Target recalculation

$$\text{next target} = \begin{cases} \frac{1}{\tau} \cdot T & \text{if } \frac{n_0}{n} \cdot T_0 < \frac{1}{\tau} \cdot T; \\ \tau \cdot T & \text{if } \frac{n_0}{n} \cdot T_0 > \tau \cdot T; \\ \frac{n_0}{n} \cdot T_0 & \text{otherwise} \end{cases}$$

- Recalculation occurs at the end of every “epoch”
 - m : epoch length in blocks (in Bitcoin: 2016)
- n_0 : estimation of number of ready parties at the system’s onset (party=CPU)
- T_0 : initial target
- τ : recalculation threshold parameter (in Bitcoin: 4)
- T : target in effect
- $n = m/(pT\Delta)$: the “effective” number of parties in the epoch
 - Δ : last epoch’s duration based on block timestamps
 - pT : probability of a single party being successful in PoW in a round

Clay pigeon shooting game



Clay pigeon shooting game

- Suppose you shoot on targets successively against an opponent
 - your success probability: 0.3
 - your opponent's success probability: 0.4
 - you shoot in sequence 1000 targets
 - winner is the one that got the most hits
- What is your probability of winning?

Chernoff Bounds

Let: $\delta > 0, \mathbf{Prob}[X_i = 1] = p_i, \mu = \sum_{i=1}^n p_i$

$$\mathbf{Prob}\left[\sum_{i=1}^n X_i \geq (1 + \delta)\mu\right] \leq \exp(-\delta^2\mu/(2 + \delta))$$

Then:

$$\mathbf{Prob}\left[\sum_{i=1}^n X_i \leq (1 - \delta)\mu\right] \leq \exp(-\delta^2\mu/2), \delta \in (0, 1)$$

Analysis, I

- You have an expectation of 300 hits
- Your opponent has an expectation of 400 hits
- What is *your* probability of winning?
 - Denote by X whether you hit a target, similarly Y for your opponent
 - From Chernoff bounds:
$$\Pr\left[\sum_{i=1}^{1000} X_i \geq 345\right] \leq \exp(-(0.15)^2 300 / 2.15) < 4.3\%$$
$$\Pr\left[\sum_{i=1}^{1000} Y_i \leq 348\right] \leq \exp(-(0.13)^2 400 / 2) < 3.5\%$$

Analysis, II

- You have an expectation of 300 hits
- Your opponent has an expectation of 400 hits
- What is *your* probability of winning?
 - Denote by X whether you hit a target, similarly Y for your opponent
 - From Chernoff bounds:
$$\Pr\left[\sum_{i=1}^{1000} X_i \geq 345\right] \leq \exp(-(0.15)^2 300 / 2.15) < 4.3\%$$
$$\Pr\left[\sum_{i=1}^{1000} Y_i \leq 348\right] \leq \exp(-(0.13)^2 400 / 2) < 3.5\%$$
- If the negation of both events happens, you will certainly lose:
 - Thus, probability of you winning is less than 8%

$$\Pr[X_{<345} \wedge Y_{>348}] = (1 - \Pr[X_{\geq 345}])(1 - \Pr[Y_{\leq 348}]) \geq 92.3\%$$

Analysis, III

- Now you are given a choice:
 - decrease the size of the clay pigeon target by a ratio β
 - augment your “kills” by multiplying with $1/\beta$
 - your accuracy is linear with β
 - your opponent will keep playing in the same way as before
- Do you accept to play like this?

Analysis, IV

- Now you are given a choice:
 - decrease the size of the clay pigeon target by a ratio β
 - augment your “kills” by multiplying with $1/\beta$
 - your accuracy is linear with β
 - your opponent will keep playing in the same way as before
- Do you accept to play like this?
- Each shot has success probability: $\Pr[X'_i = 1] = \beta \cdot \Pr[X_i = 1]$
- The score expectation of each shot remains: $E[(1/\beta)X'_i] = (1/\beta)\beta E[X_i] = E[X_i]$
- But decreasing β results in increased variance \rightarrow previous argument fails

$$\Pr\left[\sum_{i=1}^{1000} X'_i \geq 345\beta\right] \leq \exp(-(0.15)^2 300\beta / 2.15)$$

β	bound
1	$\sim 4.3\%$
0.5	$\sim 20.8\%$
0.25	$\sim 45.6\%$
0.10	$\sim 73.1\%$

The Difficulty Raising Attack

- The recalculation threshold (τ) is essential
- Without it, an adversary that has a minority of hashing power:
 - Creates a private, artificially difficult chain
 - Similar to clay pigeon shooting game, this increases the variance in its block production rate
 - Overcoming the chain of the honest parties becomes a non-negligible event

[B13] Lear Bahack. Theoretical Bitcoin Attacks with less than Half of the Computational Power (draft)